

Exercises

Exercise 3: Importing and inspecting data

Read Chapter 3 to help you complete the questions in this exercise.

1. As in previous exercises, either create a new R script or continue with your previous R script in your RStudio Project. Again, make sure you include any metadata you feel is appropriate (title, description of task, date of creation etc) and don't forget to comment out your metadata with a `#` at the beginning of the line.
2. If you haven't already, download the data file `'cardiacdata.xlsx'` from the **Data** link and save it to the **data** directory you created in Exercise 1 in your RStudio project. You will use this same dataset for the next three practicals, so it is worth putting it somewhere sensible now.
3. Notice that the file you have just downloaded is a spreadsheet, not a data file. A spreadsheet is a program's own format, and R would much rather be given plain text. Open the `'cardiacdata.xlsx'` file in Microsoft Excel (or even better use an open source equivalent - LibreOffice is a good free alternative) and save it as a tab delimited file type (see Section 3.3.1 of the Introduction to R book or watch this video if you're not sure how to do this). Name the file `'cardiacdata.txt'` and save it to the **data** directory. If you're a Windows user be careful with file extensions (things like `.txt`). By default, Windows doesn't show you the file extension (maybe the boffins at Microsoft don't think you need to know complicated things like this?) so if you enter `'cardiacdata.txt'` as a filename you might end up with a filename `'cardiacdata.txt.txt'`!
4. Time for a quick description of the `'cardiacdata.txt'` dataset to get your bearings. These data come from a cohort study of the risk factors for cardiovascular disease. A group of 163 adults, aged between 55 and 75, were examined at the start of the study and a range of measurements was taken on each of them. Each row of the dataset is one patient, identified by a unique patient number (`patno`), and each column is one of those measurements. As well as their age in years (`age`) and sex (`sex`), the study recorded blood pressure as the usual pair of numbers, systolic and diastolic, in mmHg (`systolic`, `diastolic`); body mass index in kg/m² (`bmi`); three fats measured in a blood sample, all in mmol/l, namely total cholesterol, HDL cholesterol and triglyceride (`tchol`, `hdlchol`, `triglyceride`); how many units of alcohol each person drank in the previous week (`alcohol`); and whether they were a current smoker, an ex-smoker or had never smoked (`smoking`). The structure of these data is known as a rectangular dataset (aka 'tidy' data by the cool kids). Each row is an individual observation, each column a separate variable, and the variable names are contained in the first row of the dataset (aka a header). Unlike the tidy examples in textbooks, though, this dataset is not complete. A few values are missing, and at least one value is present but cannot possibly be correct. Finding both is part of your job over this practical and the next.

5. Now let's import the `'cardiacdata.txt'` file into R. To do this you will use the workhorse function of data importing, `read.table()`. This function is incredibly flexible and can import many different file types (take a look at the help file) including our tab delimited file. Rather than relying on the default settings, get into the habit of writing out the arguments you care about every time: `header` (does the first row contain the variable names?), `sep` (what separates one value from the next?), `na.strings` (how are missing values written in the file?) and `stringsAsFactors` (should text be converted to factors?). Being explicit costs you a few seconds and saves you a great deal of confusion later. Assign the imported data to a variable with an appropriate name (such as `cardiac`). Take a look at Section 3.2.2 of the Introduction to R book or watch this video if you need any further information.

6. Once you've imported your data file nothing much seems to happen (don't worry, this is normal). To examine the contents of the dataframe one option would be to just type the variable name (`cardiac`) into the console. This is probably not a good idea and doesn't really tell you anything about the dataframe other than there's a lot of data (try it)! A slightly better option is to use the `head()` function to display the first few rows of your dataframe. Again, this is likely to just fill up your console. A better option would be to use the `names()` function which will return a vector of variable names from your dataset. However, all you get are the names of the variables but no other information. A much, much better option is to use the `str()` function to display the structure of the dataset and a neat summary of your variables. Another advantage is that you can copy this information from the console and paste it into your R script (making sure it's commented) for later reference. How many observations does this dataset have? How many variables? And now the important question: what type of variable does R think `sex` and `smoking` are, and is it right?

7. You can get another useful summary of your dataframe by using the `summary()` function. This will provide you with some useful summary statistics for each variable. Notice how the type of output depends on whether the variable is a factor or a number. Another useful feature of the `summary()` function is that it will also count the number of missing values in each variable. Which variables have missing values, and how many? While you are looking at this output, check the maximum value of `bmi`. Does that look like a plausible body mass index to you? Make a note of it and hang on to that thought until Exercise 4.

8. Summarising and manipulating dataframes is a key skill to acquire when learning R. Although there are many ways to do this, we will concentrate on using the square bracket `[ ]` notation which you used previously with vectors. The key thing to remember when using `[ ]` with dataframes is that dataframes have two dimensions (think rows and columns) so you always need to specify which rows and which columns you want inside the `[ ]` (see Section 3.4.1 and this video for some additional background information and a few examples). Let's practice.

- a) Extract all the elements of the first 10 rows and the first 4 columns of the `cardiac` dataframe and assign to a new variable called `cardiac_sub`.
- b) Next, extract all observations (remember - rows) from the `cardiac` dataframe and the columns `patno`, `sex`, `smoking` and `tchol` and assign to a variable called `cardiac_risk`.
- c) Now, extract the first 50 rows and all columns from the original dataframe and assign to a variable `cardiac_50` (there's a better way to select rows with conditional statements - you'll meet it in Exercise 4).
- d) Finally, extract all rows except the first 10 rows and all columns except the last column. Remember, for some of these questions you can specify the columns you want either by position or by name. Practice both ways. Do you have a preference? If so why?

End of Exercise 3